



Model-Based Reinforcement Learning (Day 1: Introduction)

Michael L. Littman

Rutgers University

Department of Computer Science

Rutgers Laboratory for Real-Life Reinforcement Learning




Plan

- Day 1: Introduction
 - RL
 - Q-learning
 - Convergence
 - Model-based RL
 - PAC-MDP
 - KWIK
- Day 2: Current Trends
 - Model-free RL & KWIK
 - Model/value approximation
 - Bayesian RL
 - UCT
 - Searchless planning

Start With Game...

- up
- down
- left
- right
- A
- B




Find The Ball: Elements of RL

In reinforcement learning:

- agent interacts with its environment
- perceptions (state), actions, rewards [repeat]
- task is to choose actions to maximize rewards
- complete background knowledge unavailable

Learn:

- which way to turn
- to minimize time
- to see goal (ball)
- from camera input
- given experience.





Problem To Solve

Three core issues in the dream RL system.


- generalize experience
 - use knowledge gained in similar situations
 - “learning”
- sequential decisions
 - deal properly with delayed gratification
 - “planning”
- exploration/exploitation
 - must strike a balance
 - unique to RL?



Markov Decision Processes

Model of sequential environments (Bellman 57)

- n states, k actions, discount $0 \leq \gamma \leq 1$
- step t , agent informed state is s_t , chooses a_t
- receives payoff r_t ; expected value is $R(s_t, a_t)$
- probability that next state is s' is $T(s_t, a_t, s')$




$$Q(s, a) = R(s, a) + \gamma \sum_{s'} T(s, a, s') \max_{a'} Q(s', a')$$

- Optimal behavior is $a_t = \operatorname{argmax}_a Q(s_t, a)$
- R, T unknown; some experimentation needed


Find the Ball: MDP Version

- Actions: rotate left/right
- States: orientation
- Reward: +1 for facing ball,
0 otherwise



Families of RL Approaches

policy search value-function based model based



More direct use,
less direct learning

Search for
action that
maximizes
value

Solve Bellman
equations

More direct learning,
less direct use

Q-learning

On experience $\langle s_t, a_t, r_t, s_{t+1} \rangle$:

$$\begin{aligned} Q(s_t, a_t) \leftarrow Q(s_t, a_t) \\ + \alpha_t (r_t + \gamma \max_{a'} Q(s_{t+1}, a') - Q(s_t, a_t)) \end{aligned}$$

If:

- All $\langle s, a \rangle$ visited infinitely often.
- $\sum_t \alpha_t = \infty, \sum_t \alpha_t^2 < \infty$

Then: $Q(s, a) \rightarrow Q(s, a)$ (Watkins & Dayan 92).

Model-based MDP Learner

On experience $\langle s_t, a_t, r_t, s_{t+1} \rangle$:

- $R(s_t, a_t) \leftarrow R(s_t, a_t) + \alpha_t (r_t - R(s_t, a_t))$
- $T(s_t, a_t, s_{t+1}) \leftarrow T(s_t, a_t, s_{t+1}) + \alpha_t (1 - T(s_t, a_t, s_{t+1}))$
- $T(s_t, a_t, s') \leftarrow T(s_t, a_t, s') + \alpha_t (0 - T(s_t, a_t, s'))$
- $Q(s, a) = R(s, a) + \gamma \sum_{s'} T(s, a, s') \max_{a'} Q(s', a')$

If:

- All $\langle s, a \rangle$ visited infinitely often.
- $\sum_t \alpha_t = \infty, \sum_t \alpha_t^2 < \infty$

Then: $Q(s, a) \rightarrow Q(s, a)$ (Littman 96).

PAC-MDP Reinforcement Learning



PAC: Probably approximately correct (Valiant 84)

Extended to RL (Fiechter 95, Kakade 03, etc.).


- Given $\epsilon > 0$, $\delta > 0$, k actions, n states, γ .
- We say a strategy makes a mistake each timestep t s.t. $Q(s_t, a_t) < \max_a Q(s_t, a) - \epsilon$.
- Let m be a bound on the number of mistakes that holds with probability $1 - \delta$.
- Want m poly in k , n , $1/\epsilon$, $1/\delta$, $1/(1-\gamma)$.

Must balance exploration and exploitation!

Q-learning Not PAC-MDP




- Family: initialization, exploration, α_t decay
- Combination lock



- Initialize low, random exploration (ϵ -greedy)
 - 2^n to find near-optimal reward. Keeps resetting.
 - Needs more external direction.

Model-based Can Be PAC-MDP



- Behavior differs depending on assumption

	truth: ? = low	truth: ? = high
assume: ? = low	ignore ?, optimal	ignore ?, suboptimal!
assume: ? = high	visit ?, explore	visit ?, optimal

← No PAC-MDP guarantee

← PAC-MDP if not too much exploration

Optimism Under Uncertainty

- Idea of exploration bonus well known.
- Shown to provide PAC-MDP guarantee (Kearns & Singh 02, Brafman & Tennenholtz 02).
- Key ideas:
 - Simulation lemma: Optimal for approximate model is near-optimal.
 - Explore or exploit lemma: If can't reach unknown states quickly, can achieve near-optimal reward.
- Extend to factored dynamics (Kearns & Koller 99) and metric spaces (Kakade et al. 03).



Model-free PAC-MDP

- Although not directly relevant, this problem was solved (Strehl, Li, Wiewiora, Langford, Littman 06).
- Modifies Q-learning to build rough model from recent experience.
- Total mistakes in learning $\approx nk/((1-\gamma)^8\epsilon^4)$.
- Compare to model-based methods: mistakes in learning $\approx n^2k/((1-\gamma)^6\epsilon^3)$. (Better in states, worse in horizon.)
- Lower bound, also (Li 09).



Generalization in PAC-MDP

- Can we draw on classical ML theory?
- Model learning is a supervised problem.
 - Given examples of s, a pairs, predict s' .
- Not just for table lookup anymore!
- Extend results to functions that generalize by defining the right learning problem...

3 Models for Learning Models

- PAC:** Inputs drawn from a fixed distribution. Observe inputs. For future inputs from the distribution,

Not PAC-MDP. iid assumption implies that learner cannot improve (change) behavior!
- Mistake bound:** Inputs presented online. For each, predict. If mistake, observe label. More than m mistakes,

Not PAC-MDP. Mistakes mean that a high reward can be assumed low—suboptimal.
- KWIK:** Inputs presented online. For each, can predict output or say "I don't know" and observe label. No mistakes, but can say "I don't know" m times.

Can be PAC-MDP...

KWIK Learning

- "Knows What It Knows"**
 - Like PAC, no mistakes.
 - Like mistake bound, no distribution assumption.
- Harder problem**
 - $\text{PAC} \leq \text{mistake bound} \leq \text{KWIK}$
- Very well suited to model learning:**
 - experience distribution changes during learning
 - distribution varies with behavior, which should change!
 - exploration driven by known/unknown distinction
 - don't want to be wrong and stop exploring too soon



KWIK Learn a Coin Probability

- Given m trials, x successes, $\hat{p} = x/m$
- Hoeffding bound:
 - Probability of an empirical estimate of a random variable in the range $[a,b]$ based on m samples being more than ϵ away from the true value is bounded by $\exp\left(-\frac{2m\epsilon^2}{(b-a)^2}\right)$
- So, can KWIK learn a transition probability:
 - say “I don’t know” until m is big enough so that \hat{p} is ϵ -accurate with probability $1-\delta$.



Some Things to KWIK Learn

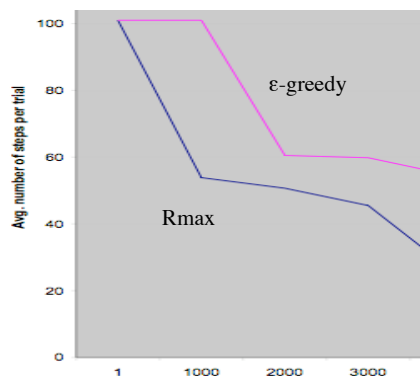
- coin probability
- an output vector, each component is KWIK learnable
 - multinomial probability (dice learning)
- a mapping from input partition to outputs where partition is known and mapping within each partition is KWIK learnable
 - That’s a standard transition function (s,a to vector of coins) (Li, Littman, Walsh 08).
- Also, union of two KWIK learnable classes.

R_{MAX} and KWIK Learning

- R_{MAX} (Brafman & Tennenholtz 02)
 - KWIK learn model ($T(s,a,\cdot)$ unknown m times).
 - For unknown parts, assume max possible reward ($Q(s,a) = r_{\max}/(1-\gamma)$).
 - Solve for Q and use resulting policy until something new becomes known.
- Total mistakes in learning $\approx n^2 k / ((1-\gamma)^3 \epsilon^3)$ (Strehl, Li, Wiewiora, Langford, Littman 06; Li 09).

R_{MAX} Speeds Learning

Task: Exit room using bird's-eye state representation.




Details: Discretized 15x15 grid x 18 orientation (4050 states);
6 actions: forward, backward, turn L, turn R, slide L, slide R.


(Nouri)

Generalizing Transitions

- In MDPs, states are viewed as independent.
 - Transition knowledge doesn't transfer.
- Real-life action outcomes generalize.
 - Learn in one state, apply to others.
- Needed:
 - MDP variants that capture transition regularities.
 - Continuous MDPs
 - RAM-MDPs
 - Factored-state MDPs
 - Object oriented MDPs



Continuous Transition Model



(Nouri)

Relocatable Action Models

Decompose MDP transitions into state-independent outcomes (Sherstov, Stone 05).

$$T'(s, a, s') = \sum_{o \text{ s.t. } \eta(s, o) = s'} t(\kappa(s), a, o)$$


- $\kappa : S \rightarrow C$ is the *type function*. It maps each state to a type (or cluster or class) $c \in C$.
- $t : C \times A \rightarrow \text{Pr}(O)$ is the *relocatable action model*. It captures the outcomes of different actions in a state-independent way by mapping a type and action to a probability distribution over possible outcomes.
- $\eta : S \times O \rightarrow S$ is the *next-state function*. It takes a state and an outcome and provides the next state that results.

RAM Example

- η : the geometry
- t : actions effects
- κ : the local walls

Example: .8 in intended direction, .1 at right angles, unless wall blocks motion.

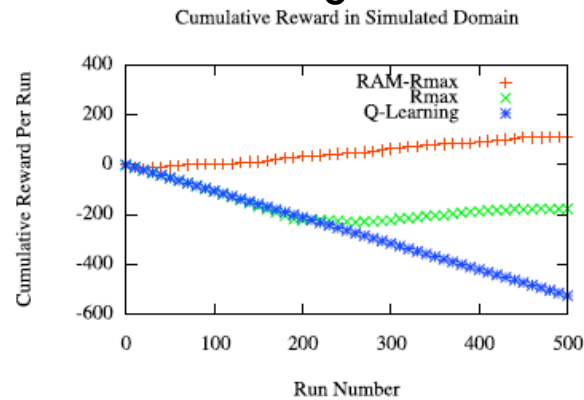
Action "go N" in a state with a walls to the N&E will go W wp .1, not move wp .9 (.8 N + .1 E).



(Leffler, Edmunds, Littman 07)



Speeds Up Learning

- Cumulative reward is larger for RAM-R_{MAX} than R_{MAX} or Q-learning.
- KWIK bound depends on classes, not states.
- (It also has more background knowledge.)



Robotic Example

- States: position and orientation
- Goal: Reach box as quickly as possible
- Types: sand, wood
- Actions: L, R, F





RAM Learning #2

QuickTime™ and a
decompressor
are needed to see this picture.

(Leffler, Mansley, Edmunds)



Factored-state MDPs

- Generalizing MDP states via DBN factoring of transition function (Boutilier et al. 99).
- 2^n states, k actions
- Blends planning-type state representations with Bayes net probability distributions
- R , T unknown; some experimentation needed
- KWIK learnable: Just a different partition of the input space.

Factored-state MDP

State is a cross product.

- Example: Taxi (**Dietterich 98**).

Primitive actions:

- N,S,E,W, pickup, dropoff.


Passenger is at R, Y, G, B.

Destination is R, Y, G, B.

Reward for successful delivery.

Approx. 500 states, but related.

- state = (taxi loc, pass. loc, dest)




Passenger: R
Destination: Y
Fuel: 5

Compact Model

- Abstraction: Use a factored (DBN) model.

independence relations
for "pickup"


(further, can use context-specific independence, **Boutilier et al. 95**)



- Model generalizes because transitions for multiple states share structure/parameters.
- If graph known, KWIK learnable:
composition of output vector and input partition.

World of Objects

- Objects in taxi:
 - taxi (location)
 - passenger (location/in taxi)
 - walls (location)
 - destination (location)



- Not *states* or state *features*, instead try *objects* and object *attributes*.
- Model: What happens when objects interact?
- More “human like” exploration.

Comparing Taxi Results

- North, not touchN(taxi, wall) → taxi.y++
- Drop, pass.in, touch(taxi, dest) → ¬pass.in
- KWIK bound: poly in types, exp in condition
- Taxi: How long until optimal behavior?

Exploration style	Algorithm	# of steps
ε greedy	Q-learning	47157
count on states	Flat Rmax	4151
count on features	Factored Rmax	1839
count on interaction	Objects	143
whatever people do	People	50

Pitfall!



A childhood dream fulfilled... (Diuk, Cohen)



Model-Based Reinforcement Learning (Day 2: Other Stuff)

Michael L. Littman

Rutgers University

Department of Computer Science

Rutgers Laboratory for Real-Life Reinforcement Learning



Plan

- Day 1: Introduction
 - RL
 - Q-learning
 - Convergence
 - Model-based RL
 - PAC-MDP
 - KWIK
- Day 2: Current Trends
 - Bayesian RL
 - Model/value approximation
 - UCT
 - Searchless planning



Structure Learning in DBNs

- Unknown structure fundamentally different.
- How can you keep statistics if you don't know what they depend on?
- Can be solved using a technique for a simpler "hidden bit" problem:
 - n -bit input, one bit (unknown) controls output
 - one output distribution if bit is on, another if off
 - Find DBN structure by same idea: one parent set controls output...

Hidden-Bit Problem

Assume the simpler deterministic setting.

Output is copy or flip of one input.

- 0110 → 0 1101 → 1 0000 → 1
- 1101 → 1 0011 → 0 1111 → 0
- 1000 → 1 1110 → 0 1100 → 1

Is it 0, 1, or “I don’t know”?

If noisy, can’t predict with each bit position separately, don’t know which to trust. Can learn about all 2^n bit patterns separately, but that’s too much.

Hidden-bit Problem via KWIK

- Can observe predictions to figure out which of k “adaptive meteorologists” to trust (Strehl, Diuk, Littman 07; Diuk et al. 09).

- Solvable with bound of $O\left(\frac{k}{\epsilon^2} \ln \frac{k}{\delta}\right) + \sum_{i=1}^k \zeta_i\left(\frac{\epsilon}{8}, \frac{\delta}{k+1}\right)$

- By considering all k -size parent sets, get a structure-learning algorithm with a KWIK bound of

$$\kappa = O\left(\frac{n^{D+3}AD}{\epsilon^3(1-\gamma)^6} \ln \frac{nA}{\delta} \ln \frac{1}{\epsilon(1-\gamma)}\right)$$

Artificial Stock Example


Discovers the structure and exploits it much faster than R_{MAX} can learn the MDP.

Factored- R_{MAX} :

Knows DBNs

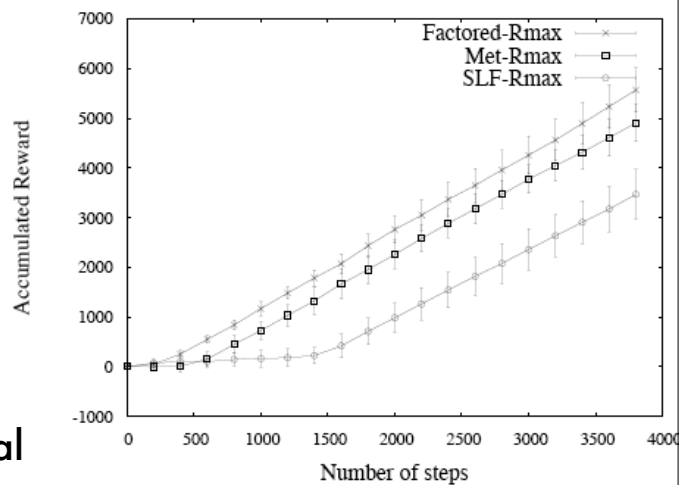
SLF- R_{MAX} : Knows size of parent sets

R_{MAX} : It's an MDP



Improved Bounds

- SLF- R_{MAX} runs in roughly $k^2 \log k$
- MET- R_{MAX} faster, like $k \log k$
- Bounds weak, but suggest a better algorithm!
- Also selected visual pattern for terrain learning.





Many Learnable Problems

Many hypothesis classes KWIK learnable:

- coin flip probability
- Dynamic Bayes net probabilities given graph
- k Dynamic Bayes net
- k Meteorologist problem
- k -CNF
- k -depth decision tree
- unions of KWIK-learnable classes
- k feature linear function



Grid World Demo (expt2)

- Unknown: What's a wall?
- (Purpose: What doesn't KWIK do?)




People Learn to Learn

- expt1: Always + (no mistakes)
- expt2: Always one shape (one mistake).
- expt3: Always some feature (two mistakes).
- Last maze is always +, but people perform differently depending on their experience.
- Transfer learning in RL (Taylor & Stone 07, e.g.).
- KWIK can learn any of these classes, but if all are possible, devolves to worst case.



Playing the Odds

- Standard PAC-MDP algorithms can't say:
 - I know you told me all states independent,
 - but every wall I've seen has been painful.
 - Can I just walk around now, please?
- Rmax vs. RAM-Rmax
 - Rmax: states independent
 - RAM-Rmax: Types known
- What if states "cluster"?
 - new state likely to be familiar





Bayesian Perspective

- Start with a prior over models.
- Maintain a posterior across tasks.
- Now, we can talk about more/less likely models instead of just *possible* models.
- How can we use the Bayesian view in exploration?



Bayes Optimal Exploration


- With a Bayesian representation of models, we can plan in the space of posteriors.
 - Can use posterior to evaluate the likelihood of any possible outcome of an action.
 - Can model how that outcome will change the posterior.
 - Can choose actions that truly maximize expected reward: No artificial distinction between exploring and exploiting or learning and acting.
- Hideously intractable except in some special cases (bandits, short horizons).

Concrete Example

- MDP has one state, 3 actions (bandit)

- $X: \{.7 \ .1 \ .8\}$, $Y: \{.8 \ .6 \ .7\}$, $\gamma = 0.8$

- Prior: $\langle .50, .50 \rangle$ ($1/2 X$, $1/2 Y$)




Concrete Example

- MDP has one state, 3 actions (bandit)


- $X: \{.7 \ .1 \ .8\}$, $Y: \{.8 \ .6 \ .7\}$, $\gamma = 0.8$

- Prior: $\langle .50, .50 \rangle$ ($1/2 X$, $1/2 Y$)



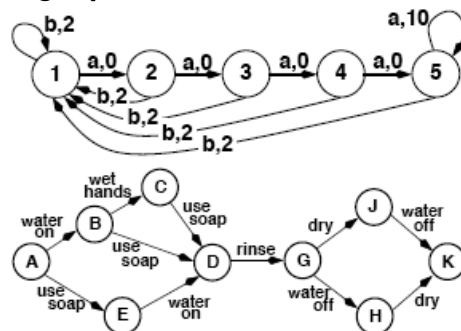
Representing Posteriors

- $T: s, a \rightarrow$ multinomial over states
- If independent for each s, a :
Dirichlet!
- Keep counts for each observed outcome.
- Can recover uncertainty in overall estimate.
- Unlike example, distribution over an infinite set.



Bayes Optimal Plans

- Many attempts (Duff & Barto 97; Dearden et al. 99)
- State of the art, BEETLE (Poupart et al. 06)
 - Latest ideas from solving continuous POMDPs
 - α functions are multivariate polynomials + PBVI
 - Can exploit “parameter tying” prior
 - Near optimal plan in “combination lock”.
 - Less optimal in bigger problem.





Near Bayes Optimal Behavior

- Recall PAC-MDP, whp makes few mistakes.
- Near Bayesian: mistakes are actions taken with values far from Bayes optimal.
- Bayesian Exploration Bonus (Kolter & Ng 09) keeps mean of posterior and adds $1/n$ bonus to actions taken n times.
 - BEB is computationally simple.
 - BEB is Near Bayesian.
 - BEB is not PAC-MDP, though...




Bayes Optimal not PAC-MDP

- Examples where Bayes optimal does not find near optimal actions (Kolter & Ng 09; Li 09)
- Not clear which is “right”.
- Who gets near optimal reward?
 - PAC-MDP: Future self
 - Near Bayesian: Current self
- Human behavior somewhere in between?
 - Hyperbolic discounting

PAC-MDP with Bayesian Priors

- With a prior that all similar colored squares are the same, we can bound the chance generalization will lead to sub-optimality.
- Idea: Don't worry about it if it's small!



$X: \{.7 \ .1 \ .8\}, Y: \{.8 \ .6 \ .7\}$

$\epsilon=0.0001, \delta=0.05$

$<.99, .01>$

R is near optimal whp

BOSS: Algorithmic Approach

- Optimism under uncertainty, not Bayes optimal
 - Sample models from the posterior.
 - Stitch together into a meta-MDP.
 - Solve to find optimal behavior: best of sampled set
 - Act accordingly until something new learned.
- If set big, near optimality whp (Asmuth et al. 09)
- Several ideas appear to be viable here

$$O\left(\frac{SAB}{\epsilon(1-\gamma)^2} \ln \frac{1}{\delta} \ln \frac{1}{\epsilon(1-\gamma)}\right)$$

BOSS in Maze

- To learn in maze:
 - Chinese Restaurant Process prior
 - Finds (empirical) clusters
 - Outperforms Rmax, 1-cluster RAM-Rmax




- Fewer than states
- Fewer than types
- Some types grouped
- Rare states nonsense

Computation Matters


- Learning/exploration can be made efficient
 - model-based RL
 - PAC-MDP for studying efficient learning
 - KWIK for acquiring transition model
- Planning “just” a computational problem.
 - But, with powerful generalization, can quickly learn accurate yet intractible models!
 - Something needs to be done or the models are useless. (Not as focused on guarantees today.)

"Nesting" RL Approaches


policy search




value function




model-based



policy search inside model-based




value function inside model-based



Example: Autonomous Flight

- Outer approach: Model-based RL.
 - Experts parameterize model space
 - Parameters learned quickly from expert demonstration (no exploration needed)
- Resulting model very high dimensional (S,A)
- Inner approach: Policy-search RL.
 - Experts parameterize space of policies
 - Offline search finds excellent policy on model
 - Methodology robust to error in model
- Learns amazing stunts (Ng et al. 03).




Tricks and Treats

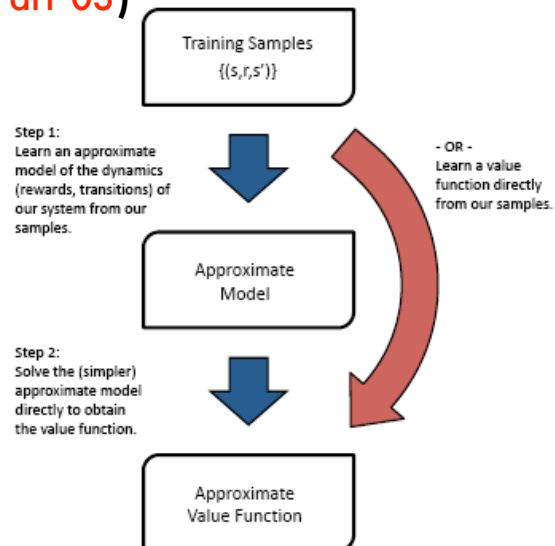


Linear Models

- Linear value function approaches: LSTD/LSPI
(Boyan 99; Lagoudakis & Parr 03)






- Give the same result!
(Parr et al. 08)




Fitted Value Iteration

- Represent value function via anchor points and local smoothing (Gordon 95)
- Some guarantees if points densely sampled (Chow & Tsitsiklis 91)
- Combined with KWIK learning of model (Brunskill et al. 08)




UCT: Upper Conf. in Trees

- Narrow, deep game-tree search via bandits (Kocsis & Szepesvári 06)
- Huge success in Go (Gelly & Wang 06)
- Good fit w/ learned model.
 - Just needs to be able to simulate transitions
 - KWIK-like methods are also “query” based
- Not aware of work using it in RL setting.




Do Kids Explore?

- Statistics of play sensitive to confounding
- Show kid 2-lever toy (Schulz/Bonawitz 07).
 - Demonstrate both. Kid becomes interested in new toy.
 - Demonstrate them together. Kids stays interested in old toy.
- Experiment design intractable. KWIK-like heuristic?



Old Toy



New Toy

Do People Explore? (xkcd)





Wrap Up

- Introduction
 - Q-learning
 - MDPs
 - Model-based RL
 - PAC-MDP
 - KWIK
- Current topics
 - Bayesian RL
 - “Recursive” RL
 - Function approximation
 - Human exploration